


# Seneca



## Computer Vision

# Simple Image Processing and Drawing Tools

Seneca Polytechnic

Credit: Vida Movahedi

# Overview

- Simple Image Processing
- Histogram Equalization
- Drawing Tools
- Mouse Callback

# Simple Image Processing

# What is Computer Vision? (Reminder)

- “The science of **creating a similar capability** [as human vision] in computers and, if possible, **to improve upon it**”
- “Computer vision is the automatic analysis of images and videos by computers in order to gain some understanding of the world ... and to **emulate the capabilities of human vision**”
- “The transformation of data from a still or video camera into either a **decision or a new representation** ... to **achieve some particular goal**”

# Digital Image Processing

- Digital image processing is the use of a digital computer to process digital images through an algorithm **to improve the quality of images**.
- It is a subcategory or a subfield of **digital signal processing**
- Many of the techniques of digital image processing started developing in the 1960s, at Bell Laboratories, the Jet Propulsion Laboratory, Massachusetts Institute of Technology, University of Maryland, and a few other research facilities
- We saw some examples last week: DSP, Image Compression, Color Models

# Image Processing

- Not the main focus in Computer Vision
- However, used as a **pre-processing step** for computer vision applications to improve the quality of images, to prepare images for AI/ML algorithms, ...
- Examples:
  - Increase brightness
  - Reduce image noise
  - Rotate image
  - Crop image
  - Resize image
  - Filtering and smoothing

# Point & Neighborhood Operators [1]

- **Point Operators:**

The value of each pixel in the **output** depends **only** on the value of the **same pixel** in the input (and possibly **some global information** or some parameters)

Example: Brightness Adjustment

- **Neighborhood Operators:**

The value of each pixel in the **output** depends on the value of the **pixel and** the value of **its neighbors** in the input

Example: Smoothing, Blurring, Filtering

# Point Operators [1]

## Point Operators:

The value of each pixel in the **output** depends **only** on the value of the **same pixel** in the input (and possibly **some global information** or some parameters)

Example: brightness adjustment

Example: Assuming **one color channel** for simplicity

- $I_{\text{out}}(i, j) = f(I_{\text{in}}(i, j))$ , for some function  $f$



# Point Operators [1]

## Brightness Adjustment

- **Addition** with a constant, positive or negative – (Brightness changes the **overall** lightness of the image)

$$I_{\text{out}}(i, j) = I_{\text{in}}(i, j) + b$$

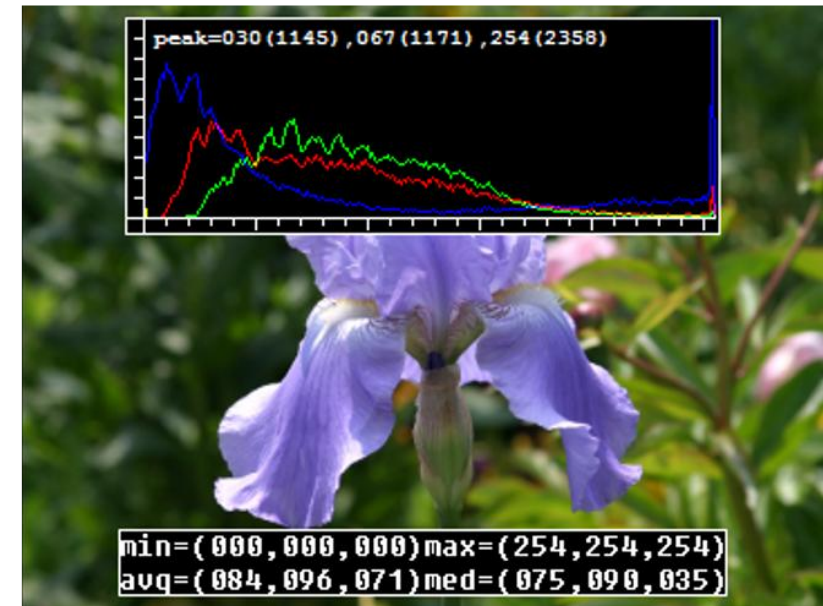
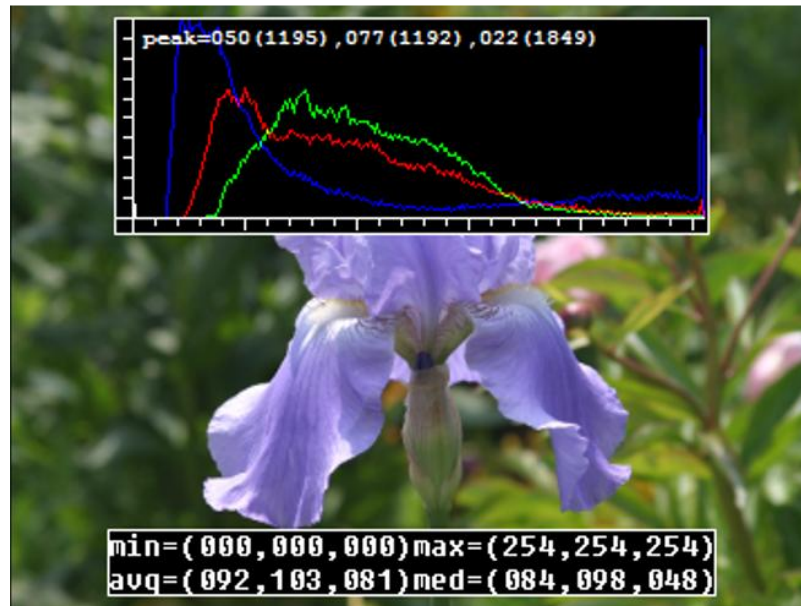
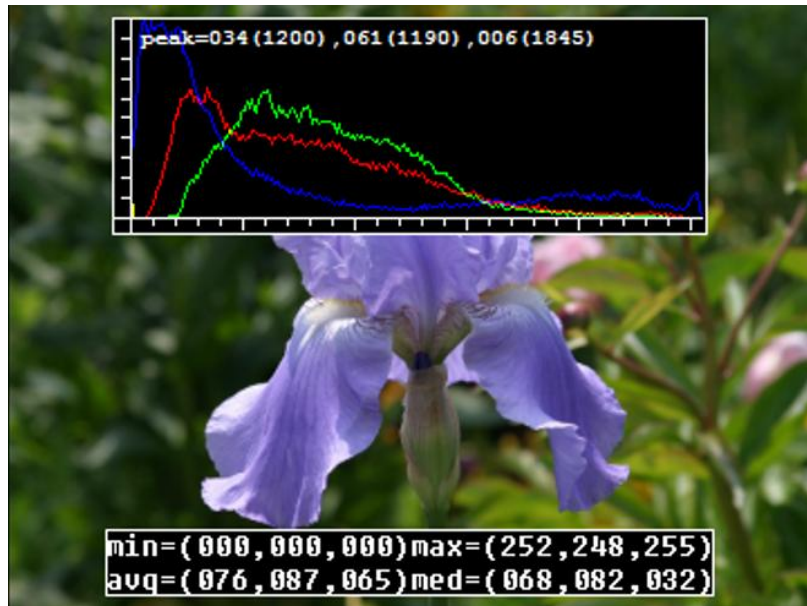
## Contrast Adjustment

- **Multiplication** with a constant, smaller than 1 or greater than 1 – (Contrast adjusts the difference between the darkest and lightest colors)

$$I_{\text{out}}(i, j) = aI_{\text{in}}(i, j)$$

# Contrast & Brightness Adjustment [1]

- Left: original image with its three color (per-channel) histograms;
- Middle: brightness increased (additive offset,  $b = 16$ );
- Right: contrast increased (multiplicative gain,  $a = 1.1$ );



# Addition & Subtraction (Brightness Adjustment)

- Using arithmetic operations in Numpy (does not work when values can go above 255 or lower than 0 )
- Use **OpenCV Functions** instead

```
// Make brighter
```

```
img2 = cv.add(img,  
              np.ones(img.shape, dtype = "uint8") * 50)
```

```
// Make darker
```

```
img3 = cv.subtract(img,  
                  np.ones(img.shape, dtype = "uint8") * 100)
```






# Multiply & Divide (Contrast Adjustment )

```
// Make bright colors brighter- use scale > 1  
img2 = cv.multiply(img,  
                    np.ones(img.shape, dtype = "uint8"),  
                    scale = 1.4)
```

```
// Make dark colors darker- use scale < 1  
img3 = cv.multiply(img,  
                    np.ones(img.shape, dtype = "uint8"),  
                    scale = 0.6)
```



# Linear Blend (Weighted Image Addition)

- Two input images, *img1* and *img2*
- $dst = \alpha \cdot img1 + \beta \cdot img2 + \gamma$

```
img1 = cv.imread("Trillium.jpg")  
img2 = cv.imread("flower.jpg")  
img2 = cv.resize(img2, (img1.shape[1], img1.shape[0]))  
img3 = cv.addWeighted(img1, 0.6, img2, 0.4, 0)
```

$\alpha = 1, \beta = 0$



$\alpha = 0.6, \beta = 0.4$



$\alpha = 0.4, \beta = 0.6$




$\alpha = 0, \beta = 1$



# Histogram Equalization

# Histogram

- A **visual representation** of an image where the frequency of each image value (brightness/intensity) is determined
- The **count** of the number of pixels at each value (**frequency** of each intensity value)
- X axis: **Intensity** values (e.g. 0 to 255)
- Y axis: **Number of Pixels** in the image having that value



A grey-scale image and the histogram derived from it [3]



- The algorithm to derive such a **histogram  $h(g)$**  from a grey-scale image  $f(i,j)$  is:

### Algorithm 3.1

---

```
// Initialise the histogram
for (g = 0; g <= 255; g++)
    h(g) = 0
// Compute the histogram
for (i = 0; i < MAXcolumn; i++)
    for (j = 0; j < MAXrow; j++)
        h(f(i,j))++
```

---

[3]

- The above code assumes 256 bins (0 to 255), and counts the number of pixels with the corresponding value




# Histogram

- Histogram contains **global information** about the image and that information is completely independent of the **position** and the **orientation** of objects in the scene
- The number of **bins** is often set to smaller numbers, for example 4
- **Normalized histogram**: Divide all numbers by the total number of pixels in the image, therefore


$$\Rightarrow \sum_{g=0}^{255} h(g) = 1$$

# Color Histograms


One Histogram per  
Color Map or Each  
Channel




Red



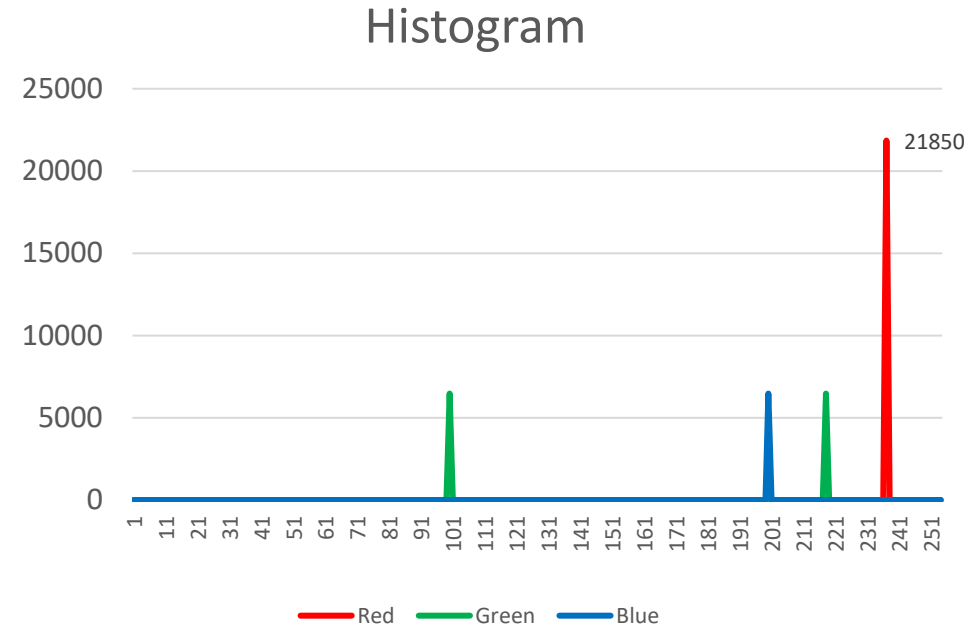

Green



Blue



# Color Histograms





**Overlapping** the three histograms on top of each other for **comparing** the **size** of the objects with specific colors

(This histogram **does not include** counts for values 0 and 255)

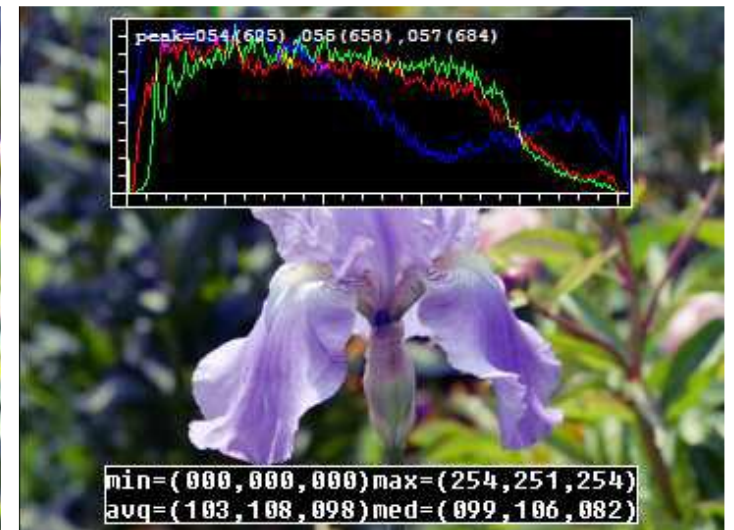

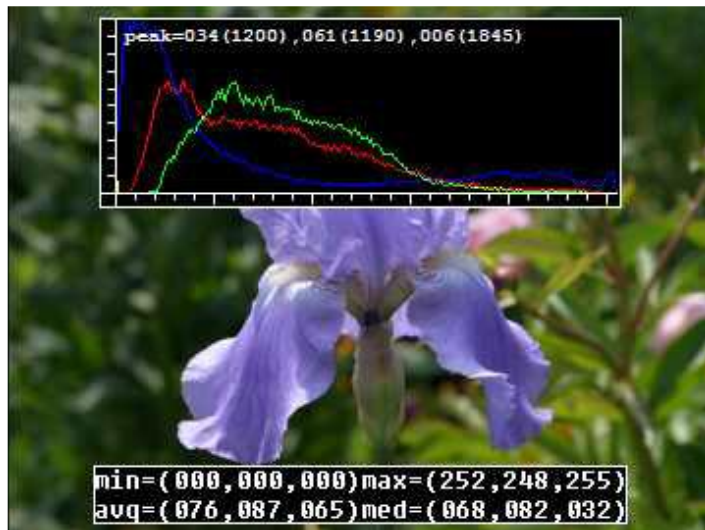
# Histogram Equalization [2]

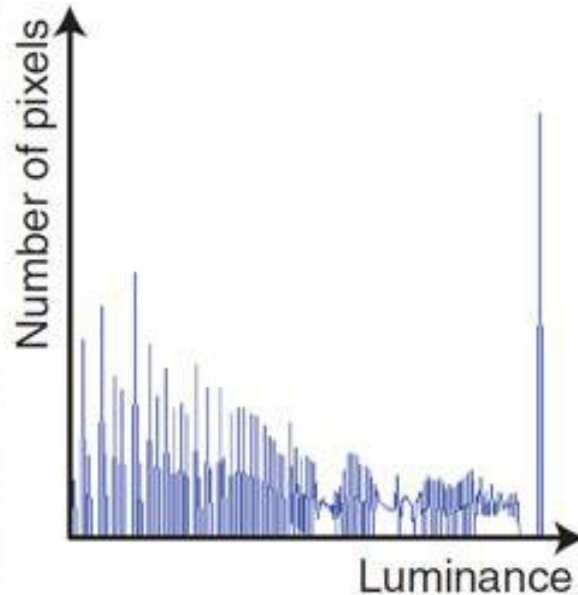
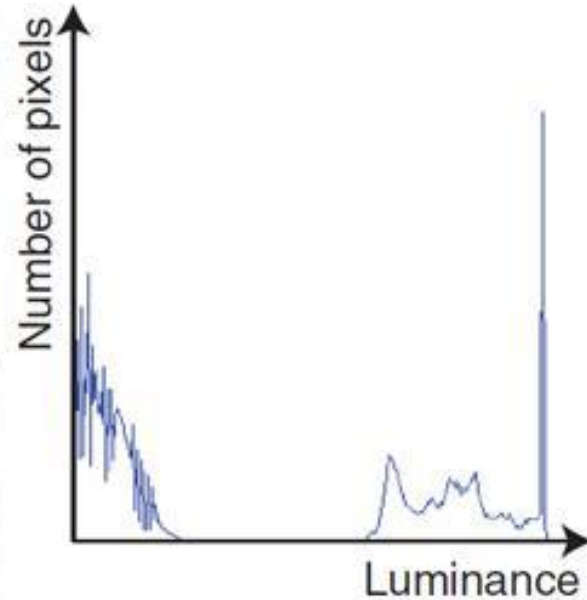
- When there is an **excess of pixels** in a certain range, we use Histogram Equalization.
- A technique for improving the distribution of grey-scales in an image
- Points of some common grey-scales in the input are mapped to **multiple different grey-scales** in the output



# Histogram Equalization [1]

- Mapping the values in a way that results in a **flatter distribution** → The trick to finding such intensity mapping function is the same one to generate random samples from a probability density function.





- When equalizing a color image, we generally only equalize the luminance channel as otherwise the colors can become distorted.
- Already implemented in OpenCV (`cv::equalizeHist`)[2] for **single channel image**
- Example: A color image and its luminance histogram (top), together with a histogram equalized version of the image and the resultant luminance histogram (bottom)[3]

# Beyond Histogram Equalization

- Histogram is fast and easy to compute.
- Size can easily be normalized so that different image histograms can be compared → The normalized histogram gives the percentage of pixels that have gray tone  $i$ .
- Histogram can be used to match color histograms for database query or classification.



# Beyond Histogram Equalization



Retrieving images that are similar to a given image or that contain a particular content (Content Based Image Retrieval – CBIR).

- It is possible to analyze the color distribution in an image by comparing histograms derived from the images.
  - Compare all pixels? Computationally expensive
  - Compare their histograms first; Calculate the histogram difference or distance
- Image search engines do this mostly by using meta-data tags associated with each image





# Beyond Histogram Equalization

- Finding the most similar image in the database involves identifying the 3D color histogram that best matches the 3D color histogram of the query image.
- In multidimensional histograms the bins are not 1-dimensional but e.g., 3-dimensional for 3-channel images

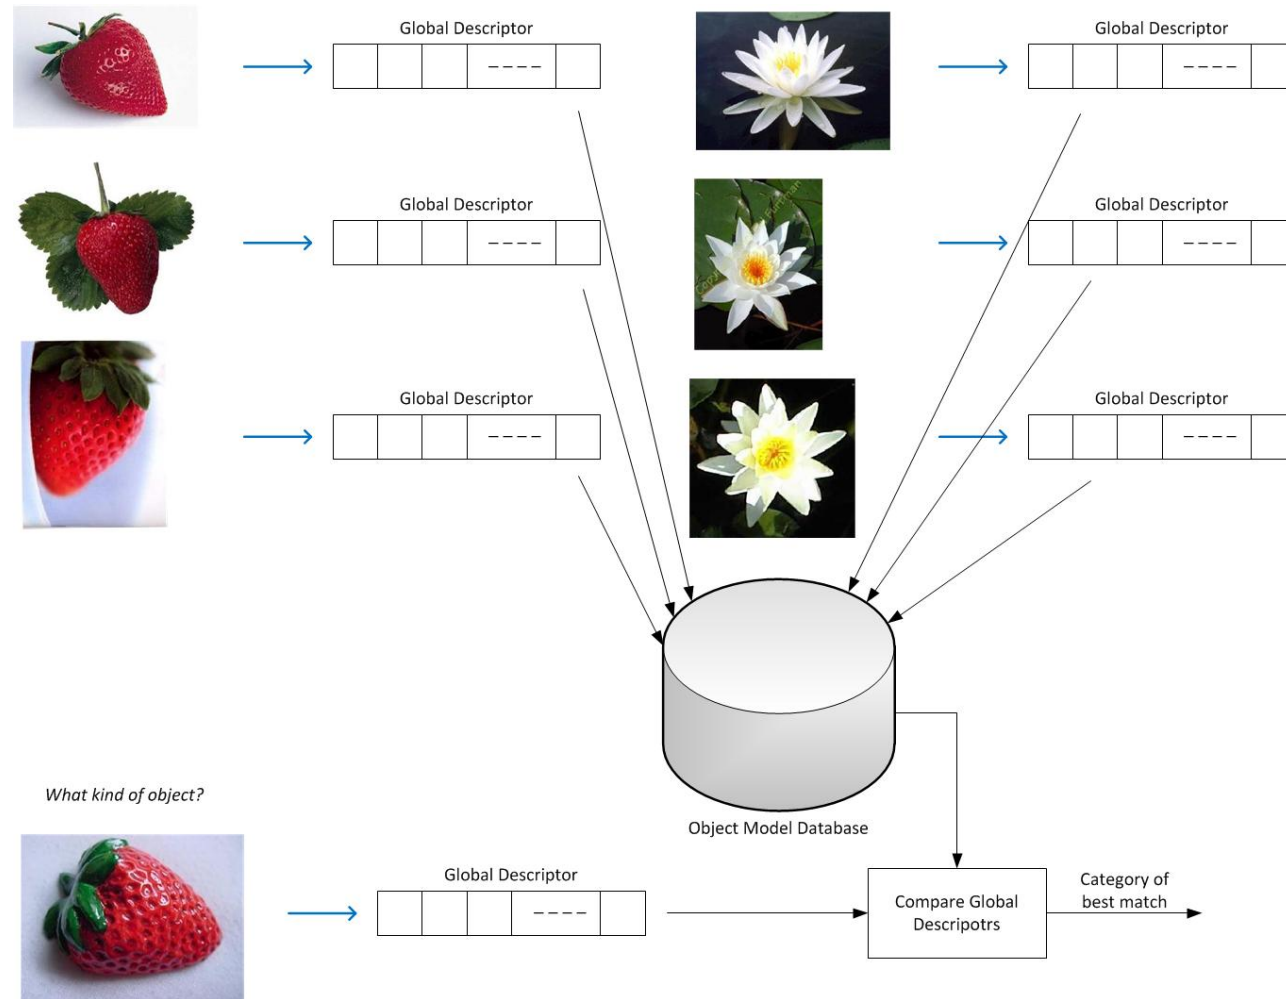


# Beyond Histogram Equalization

- For an 8-bin histogram, the 3D histogram will have 512 bins
  - The bin  $[0,0,0]$  counts the number of pixels in the image, for which the values of the red, green and blue channel are in the range between 0 and 15 (inclusive).
  - The bin  $[7,7,7]$  counts the number of pixels in the image, for which the values of the red, green and blue channel are in the range between 240 and 255 (inclusive).



# Beyond Histogram Equalization



# Histogram Comparison [3]

Compare their histograms → Calculate the histogram similarity or difference/distance

- $D_1 = \sum_i (h_1(i) - h_2(i))$

Other **distance/similarity functions** for comparing histograms exist that perform better than the above function such as correlation, Chi-Square, intersection, Earth Mover's Distance, ...



Best matching images (based on histogram comparison). The reference image is the one on the top left (which gets a perfect matching score). The metrics shown in red are the **correlation** scores [3]

# Summary

- Simple image processing can be done by image arithmetic, such as adding, subtracting, multiplying, or dividing by a constant, or adding or subtracting two images.
- A histogram is a visual representation of the count of pixels (frequency) at each intensity or color value. Histogram equalization is a technique for mapping values to a flatter distribution and often results in an improvement to the image.
- Two types of connectivity in an image are: 4-connectivity and 8-connectivity.
- OpenCV has many tools and functions for drawing shapes on an image and listening to mouse events.



# References

[1] **Computer Vision: Algorithms and Applications** by R. Szeliski

[Computer Vision: Algorithms and Applications, 2nd ed. \(szeliski.org\)](http://szeliski.org)

[2] **Learning OpenCV 3** by A. Kaehler & G. Bradski

Available online via Seneca Libraries: [Learning OpenCV 3 : computer Vision in C++ with the OpenCV Library - Seneca \(exlibrisgroup.com\)](http://exlibrisgroup.com)

[3] **A Practical Introduction to Computer Vision with OpenCV** by Kenneth Dawson-Howe

Available online via Seneca Libraries: [A Practical Introduction to Computer Vision with OpenCV. - Seneca \(exlibrisgroup.com\)](http://exlibrisgroup.com)

# Readings

# Chapter 3 [1]

# Chapters 6, 7, 9 & 13 [3]

- **[1] A Practical Introduction to Computer Vision with OpenCV**  
by Kenneth Dawson-Howe  
Available online via Seneca Libraries: [A Practical Introduction to Computer Vision with OpenCV. - Seneca \(exlibrisgroup.com\)](https://exlibrisgroup.com/10.1017/9781009054444)
- **[2] Learning OpenCV 4 Computer Vision with Python 3**  
by J. Howse & J. Minichino  
Available online via Seneca Libraries: [Learning OpenCV 4 Computer Vision with Python 3 : get to grips with tools, techniques, and algorithms for computer vision and machine learning - Seneca \(exlibrisgroup.com\)](https://exlibrisgroup.com/10.1017/9781009054444)
- **[3] Learning OpenCV 3**  
by A. Kaehler & G. Bradski  
Available online via Seneca Libraries: [Learning OpenCV 3 : computer Vision in C++ with the OpenCV Library - Seneca \(exlibrisgroup.com\)](https://exlibrisgroup.com/10.1017/9781009054444)



# OpenCV: Drawing Tools

# Draw a Simple Line


```
import numpy as np  
import cv2 as cv
```

```
# Create a black image
```

```
img = np.zeros((512,512,3), np.uint8)
```

```
# Draw a diagonal blue line with thickness of 5 px
```

```
cv.line(img, (0,0), (511,511), (255,0,0), 5)
```




# Draw a Simple Line

```
cv.line(  
    img,          # Image to be drawn  
    pt1,          # Starting point  
    pt2,          # Endpoint  
    color[,       # Color BGR form  
    thickness[,   # Thickness of line  
    lineType[,    # Connectedness, 4, 8, or cv.LINE_AA  
    shift]])      # Bits of radius to treat as fraction  
  
    ) ->    img
```



# Draw a Simple Rectangle

- `cv.rectangle(img, (384, 0), (510, 128), (0, 255, 0), 3)`



# Drawing a circle

- `cv.circle(img, (447, 63), 63, (0, 0, 255), 5)`
- `cv.circle(img, (447, 63), 63, (0, 0, 255), cv.FILLED)`



# More Drawing Functions

- See [OpenCV: Drawing Functions in OpenCV](#)

```
cv.ellipse(img,(256,256),(100,50),0,0,180,(255,0,0),-1)
```


```
pts = np.array([[10,5],[20,30],[70,20],[50,10]], np.int32)
```

```
pts = pts.reshape((-1,1,2))
```

```
cv.polylines(img,[pts],True,(0,255,255))
```

```
font = cv.FONT_HERSHEY_SIMPLEX
```

```
cv.putText(img,'OpenCV',(10,500), font, 4,(255,255,255),2,cv.LINE_AA)
```



# Draw with Your Mouse

[OpenCV: Drawing Functions in OpenCV](#)

# Draw a Circle When Double Click

```
import numpy as np
import cv2 as cv

# mouse callback function
def draw_circle(event, x, y, flags, param):
    if event == cv.EVENT_LBUTTONDBLCLK:
        cv.circle(img,(x,y),100,(255,0,0),-1)

# Create a black image, a window and bind the function to window
img = np.zeros((512,512,3), np.uint8)
cv.namedWindow('image')
cv.setMouseCallback('image',draw_circle)

while(1):
    cv.imshow('image',img)
    if cv.waitKey(20) & 0xFF == 27:
        break
cv.destroyAllWindows()
```



# Listen to Mouse Events

## 1. Define a callback function

- A function that OpenCV calls whenever a mouse events occurs
- For example, when the left button is pressed (down) or released (up), when the mouse is moved



```
void your_mouse_callback(  
    int    event,           // Event type (see Table 9-1)  
    int    x,              // x-location of mouse event  
    int    y,              // y-location of mouse event  
    int    flags,          // More details on event (see Table 9-2)  
    void*  param           // Parameters from cv::setMouseCallback()  
);
```

## 2. Register the callback function for the window

```
void cv::setMouseCallback(  
    const string&    windowName,           // Handle used to identify window  
    cv::MouseCallback your_mouse_callback, // Callback function  
    void*            param = NULL         // Additional parameters for callback fn.  
);
```


# Connectivity

- 4- connectivity: A pixel P is considered connected to 4 neighbors
- 8-connectivity: A pixel P is considered connected to 8 neighbors




# Drawing on images [2]


- `lineType`
  - Integer (4, 8 (default) , or `cv::LINE_AA`)
  - 8 is the default for `lineType`



(a) 4-connected Bresenham  
Algorithm



(b) 8-connected Bresenham  
Algorithm



(c) Anti-aliased Line With  
Gaussian Smoothing